
Towards Better Generalization via Distributional Input Projection Network

Yifan Hao^{1*} Yanxin Lu^{1*} Xinwei Shen² Tong Zhang¹

¹University of Illinois Urbana-Champaign

²Eidgenössische Technische Hochschule Zurich Departement Mathematik

Abstract

As overparameterized models become increasingly prevalent, training loss alone offers limited insight into generalization performance. While smoothness has been linked to improved generalization across various settings, directly enforcing smoothness in neural networks remains challenging. To address this, we introduce Distributional Input Projection Networks (DIPNet), a novel framework that projects inputs into learnable distributions at each layer. This distributional representation induces a smoother loss landscape with respect to the input, promoting better generalization. We provide theoretical analysis showing that DIPNet reduces both local smoothness measures and the Lipschitz constant of the network, contributing to improved generalization performance. Empirically, we validate DIPNet across a wide range of architectures and tasks, including Vision Transformers (ViTs), Large Language Models (LLMs), ResNet and MLPs. Our method consistently enhances test performance under standard settings, adversarial attacks, out-of-distribution inputs, and reasoning benchmarks. We demonstrate that the proposed input projection strategy can be seamlessly integrated into existing models, providing a general and effective approach for boosting generalization performance in modern deep learning.

1 Introduction

Overparameterization has become a defining feature of modern deep learning. From vision transformers to large language models, today’s networks often possess far more parameters than training examples. While this overparameterization enables remarkable expressivity and optimization ease, it is of great importance to identify models with strong generalization performance (i.e, the excellent performance on unseen test data).

Recent work has sought to characterize generalization through various geometric and functional properties of the learned models. Notably, Johnson and Zhang (2023) argue that the generalization gap can be largely attributed to two components: inconsistency and instability, with the latter being closely related to the Lipschitz continuity of the learned function. In parallel, another prominent line of research focuses on sharpness, typically quantified via the spectral norm of the Hessian. Building on this perspective, Foret et al. (2020) proposed Sharpness-Aware Minimization (SAM), a technique aimed at locating flatter regions of the loss landscape. SAM has been successfully applied in diverse domains including computer vision (Chen et al., 2021), natural language processing (Bahri et al., 2021), and bi-level optimization (Abbas et al., 2022).

However, enforcing low sharpness, or low Lipschitz constants during training remains a significant challenge—particularly without sacrificing generalization. For instance, there is no guarantee that a

*Random order, equal contribution.

Emails: yifanh12@illinois.edu, yanxin14@illinois.edu, xinwei.shen@stat.math.ethz.ch, tongzhang@tongzhang-ml.org




Figure 1: Pipeline of Distributional Input Projection Network.

model with low Lipschitz on the training distribution will remain a low Lipschitz on the test distribution. Adversarial training (Madry et al., 2017) and Random Smoothing (Cohen et al., 2019) have proven effective in reducing the Lipschitz constant and improving robustness, but often introduces a trade-off between robustness and standard generalization performance (Tsipras et al., 2018; Zhang et al., 2019; Javanmard et al., 2020; Donhauser et al., 2021; Dobriban et al., 2023; Hao and Zhang, 2024). Likewise, while SAM can find flatter regions of the landscape, it does not always yield true minimizers, even in sufficiently flat neighborhoods (Yue et al., 2023).

An increasingly promising direction involves promoting smoothness in the learned models. Some recent works (Johnson and Zhang, 2023; Shen and Meinshausen, 2023) suggest that smoother input-output mappings lead to improved robustness and generalization. Smoothness mitigates the model’s sensitivity to small input perturbations, thereby enhancing stability across data distributions. Nonetheless, directly enforcing smoothness in deep networks remains difficult due to their nonlinear and high-dimensional nature.

In this paper, we introduce Distributional Input Projection Networks (DIPNet)—a novel architectural framework that improves generalization by implicitly promoting smoothness. Rather than passing deterministic inputs through the network, DIPNet projects inputs into learnable distributions at each layer, enabling the model to reason over localized neighborhoods in the input space. This structured stochasticity acts as a form of regularization, smoothing the loss landscape and mitigating sensitivity to input variation. As shown in Figure 1, the proposed framework consistently improves generalization across a wide range of datasets and model architectures, with particularly notable gains under adversarial attacks and distribution shifts.

Moreover, DIPNet is broadly applicable and can be integrated into existing architectures with minimal overhead. Its distributional nature also draws inspiration from variational inference, which motivates our efficient training procedure with a stability-promoting penalty.

Our main contributions are as follows:

1. We introduce DIPNet, a novel architectural framework that enhances smoothness by projecting inputs into learnable distributions at each layer of the network.
2. We develop an efficient training method for DIPNet, inspired by variational inference, which also incorporates a stability penalty to promote robustness and regularization.
3. We provide theoretical guarantees showing that DIPNet reduces function smoothness measures and Lipschitz constants, offering insight into its generalization benefits.
4. We conduct extensive experiments across a diverse set of tasks and architectures—including ViTs, LLMs, MLPs, ResNets, adversarial robustness settings, out-of-distribution (OOD) generalization, and reasoning benchmarks—demonstrating consistent improvements in test-time performance. Such improvements are particularly significant under adversarial attacks or distribution shifts.

2 Related work

Influential factors on generalization performance. In real-world scenarios, over-parameterized neural networks often find solutions that perform optimally or near-optimally on training data, but these solutions do not always generalize well to test data. Given the significant challenges associated with generalization performance in deep neural networks, many studies have aimed to identify the key factors influencing generalization. From the parameter perspective, a considerable body of research has focused on the relationship between function sharpness and generalization (Hochreiter and Schmidhuber, 1997; Keskar et al., 2016; Izmailov et al., 2018; Jiang et al., 2019; Foret et al., 2020). For instance, Foret et al. (2020) introduced a method named *sharpness-aware-minimization* (SAM), which could reduce the sharpness of output model and subsequently improves generalization. However, even if the solution of SAM always aligns with flat loss landscape, it may not be optimal (Yue et al., 2023). From the input covariates perspective, aligning with the conventional smoothness viewpoint, Johnson and Zhang (2023) recently suggested different indicators to measure generalization, i.e., “inconsistency” (Nakkiran and Bansal, 2020; Jiang et al., 2021; Kirsch and Gal, 2022) and “instability” (Bousquet and Elisseeff, 2002; Shalev-Shwartz et al., 2010), with the latter being related to the Lipschitz norm of the output models. While adversarial training (Madry et al., 2017; Cohen et al., 2019; Salman et al., 2019; Lecuyer et al., 2019; Gowal et al., 2020; Wang et al., 2021; Zou et al., 2021) can control the Lipschitz norm, it often involves a trade-off between generalization performance and adversarial robustness (Lipschitz) (Tsipras et al., 2018; Zhang et al., 2019; Javanmard et al., 2020; Donhauser et al., 2021; Dobriban et al., 2023; Hao and Zhang, 2024). In contrast, our proposed method enhances generalization performance by projecting the input into a learnable distribution, thereby avoiding such trade-off typically observed with other approaches.

Smoothing methods. Prior work has explored enhancing smoothness through perturbation-based sampling in various contexts. For example, Shen and Meinshausen (2023) proposed Engression, a method that learns distributions via pre-additive perturbations. Similarly, diffusion models (Ho et al., 2020; Song et al., 2020; Dhariwal and Nichol, 2021; Saharia et al., 2022; Rombach et al., 2022) can be interpreted as techniques for modeling sample distributions from Gaussian noise. Another related line of research focuses on improving model robustness. In particular, certified robustness methods (Cohen et al., 2019; Salman et al., 2019; Lecuyer et al., 2019; Yang et al., 2020) aim to ensure stability under perturbations by injecting noise from specific distributions during training. Gaussian noise injection has also been widely used in data augmentation (Moreno-Barea et al., 2018) to enhance generalization by artificially expanding the training set. While our method also seeks to promote smoothness, it differs in two key ways: (i) Rather than treating noise injection as a training-only technique, we incorporate distributional projection directly into the model architecture, applying it consistently during both training and inference. This architectural integration enables theoretical guarantees for improved generalization. (ii) Instead of limiting perturbations to the input layer, we project the input into a learnable distribution at every layer of the network. This layerwise control provides enhanced stability and helps mitigate issues such as gradient explosion during training.

3 Problem Formulation and Variational Inference

In this section, we formally define the problem. Our goal is to learn a model parameterized by θ that minimizes the negative log-likelihood loss for prediction tasks:

$$\mathcal{L}(\theta) = -\mathbb{E}_x \ln \mathbb{P}(y|x, \theta).$$

Derivation of distributional input. Inspired by Shen and Meinshausen (2023), on each layer of the model, we consider to project input $x \in \mathbb{R}^p$ into a distribution $\mathcal{N}(x, \Sigma)$, where Σ is the learnable variance, to take a better prediction. With a training set containing n samples $\{x_i, y_i\}_{i=1}^n$, we can assume there exists an unobserved variable $\eta \sim \mathcal{N}(0, \gamma I_p)$, such that

$$\mathbb{P}(y|x, \eta, \theta) = \mathbb{P}(y|x + \eta, \theta). \quad (1)$$

Following the standard variational-inference derivation, we obtain the bound for $\mathcal{L}(\theta)$ as:

$$\begin{aligned}\mathcal{L}(\theta) &= -\sum_{i=1}^n \ln \mathbb{P}(y_i|x_i, \theta) = -\sum_{i=1}^n \ln \mathbb{E}_\eta \mathbb{P}(y_i, \eta|x_i, \theta) \\ &= -\sum_{i=1}^n \ln \mathbb{E}_\eta (\mathbb{P}(y_i|x_i + \eta, \theta) \mathbb{P}(\eta)) \\ &\leq -\sum_{i=1}^n \mathbb{E}_{\eta \sim q(\eta)} [\ln \mathbb{P}(y_i|x_i + \eta, \theta) + \ln \mathbb{P}(\eta) - \ln q(\eta)],\end{aligned}$$

for any distribution $q(\eta)$, where the third equality is from Eq. (1), and the last inequality is from ELBO lower bound². For simplicity, we constrain $q(\eta)$ to be a Gaussian distribution:

$$q(\eta) := \mathcal{N}(0, \Sigma), \quad \text{where} \quad \Sigma = \text{diag}\{\lambda_1, \dots, \lambda_p\},$$

then the lower bound should be:

$$\begin{aligned}& -\sum_{i=1}^n \mathbb{E}_{\eta \sim \mathcal{N}(0, \Sigma)} \left[\ln \mathbb{P}(y_i|x_i + \eta, \theta) - \frac{p \ln \gamma}{2} + \frac{\ln |\Sigma|}{2} - \frac{\eta^T \eta}{2\gamma} + \frac{\eta^T \Sigma^{-1} \eta}{2} \right] \\ &= -\sum_{i=1}^n \mathbb{E}_{\eta \sim \mathcal{N}(0, \Sigma)} \left[\ln \mathbb{P}(y_i|x_i + \eta, \theta) - \frac{p \ln \gamma}{2} + \frac{\sum_{j=1}^p \ln \lambda_j}{2} - \frac{\sum_{j=1}^p \lambda_j}{2\gamma} + \frac{p}{2} \right].\end{aligned}$$

Based on this formulation, to achieve accurate predictions, we minimize the following loss function:

$$\min_{\theta, \Sigma} \left\{ -\sum_{i=1}^n \mathbb{E}_{\eta \sim \mathcal{N}(0, \Sigma)} \ln \mathbb{P}(y_i|x_i + \eta, \theta) + \alpha \sum_{j=1}^p \lambda_j - \beta \sum_{j=1}^p \ln \lambda_j \right\}, \quad (2)$$

with tuning parameter $\alpha, \beta > 0$.

Remark 1. The term $\alpha \sum_{j=1}^p \lambda_j - \beta \sum_{j=1}^p \ln \lambda_j$ can be regarded as a special penalty term to prevent the corresponding parameters $\{\lambda_1, \dots, \lambda_p\}$ shrinking to zero during training process.

4 DIPNet: Distributional Input Projection Network

Building on the framework above, we propose a new algorithm—Distributional Input Projection Network (DIPNet), which projects the input into a learnable distribution at each layer of the neural network. An overview of the DIPNet pipeline is presented in Figure 1. We describe the key components and implementation details in this section.

4.1 Learning Distributional Input Layerwise to Minimize Loss Function

Motivated by enhancing generalization performance, DIPNet can switch the standard multi-layer neural network into a distributional input projection framework. To be specific, on each layer of the model, denoting the output of the previous layer as v , DIPNet performs the following steps:

1. Project v into a learnable Gaussian distribution $\mathcal{N}(v, \Sigma)$;
2. Sample a particle u from $\mathcal{N}(v, \Sigma)$;
3. Use u as the input to the current layer and compute its output.

Since each layer introduces stochasticity through distributional projection, the final output becomes inherently random. To ensure stable and reliable predictions, this input-output procedure is repeated k times. The final output is then obtained by distilling the results from all k sampled trajectories.

²Here is an upper bound as we consider the negative log-likelihood function.

Adding a stability penalty. While DIPNet can enforce smoothness by distributional projection, it induces instability on model output $f(x, \eta_1, \dots, \eta_L, \theta)$ (L is the number of model layers). To address this issue, we propose to penalizing the variance of the model output. To be specific, based on Eq. (2), we now formulate the loss function as:

$$\min_{\theta, \{\Sigma_l\}_{l=1}^L} \mathcal{L}_{\alpha, \beta, \lambda}(\theta, \Sigma_1, \dots, \Sigma_L) := \min_{\theta, \{\Sigma_l\}_{l=1}^L} \left\{ - \sum_{i=1}^n \mathbb{E}_{\eta_1, \dots, \eta_L} \ln \mathbb{P}(y_i | x_i, \eta_1, \dots, \eta_L, \theta) + \alpha \sum_{l=1}^L \sum_{j=1}^{p_l} \lambda_j^l - \beta \sum_{l=1}^L \sum_{j=1}^{p_l} \ln \lambda_j^l + \underbrace{\lambda \sum_{i=1}^n \mathbb{V}_{\eta_1, \dots, \eta_L} [f(x_i, \eta_1, \dots, \eta_L, \theta)]}_{\text{stability penalty}} \right\}, \quad (3)$$

where λ is a regularization parameter, and for each $l \in [L]$, we have $\eta_l \sim \mathcal{N}(0, \Sigma_l)$, and $\Sigma_l = \text{diag}\{\lambda_1^l, \dots, \lambda_{p_l}^l\}$. With a proper choice of λ , adding such penalty in training process can avoid extremely instable model output, thereby benefits generalization (Johnson and Zhang, 2023).

Remark 2. If we consider a simple one-layer linear model, i.e.,

$$\mathbb{P}(y|x + \eta, \theta) \sim \mathcal{N}(\theta^T(x + \eta), \sigma^2),$$

with a constant $\sigma > 0$. The stability penalty can be expressed as

$$\lambda \sum_{i=1}^n \mathbb{V}_{\eta \sim \mathcal{N}(0, \Sigma)} [\theta^T(x_i + \eta)] = \lambda n \theta^T \Sigma \theta,$$

which is equivalent to a ℓ_2 -norm penalty on model parameter.

Unbiased estimation on loss function. Before introducing the training algorithm formally, we first formulate the true loss function we use in practice:

$$\begin{aligned} & -\frac{1}{m} \sum_{i=1}^n \sum_{j=1}^m \ln \mathbb{P}(y_i | x_i, \{\eta_{l,i,j}\}_{l=1}^L, \theta) + \alpha \sum_{l=1}^L \sum_{j=1}^{p_l} \lambda_j^l - \beta \sum_{l=1}^L \sum_{j=1}^{p_l} \ln \lambda_j^l \\ & + \frac{\lambda}{m(m-1)} \sum_{i=1}^n \sum_{1 \leq j_1 < j_2 \leq m} \|f(x_i, \{\eta_{l,i,j_1}\}_{l=1}^L, \theta) - f(x_i, \{\eta_{l,i,j_2}\}_{l=1}^L, \theta)\|_2^2, \end{aligned} \quad (4)$$

where $\{\eta_{l,i,j}\}$ are i.i.d. sampled from $\mathcal{N}(0, \Sigma_l)$, and m is the number of sample times. This leads to the *unbiased estimator* for Eq. (3). We now formally introduce the Distributional Input Projection Network (DIPNet) algorithm, which is summarized in Algorithm 1.

Algorithm 1 Distributional Input Projection Network (DIPNet)

- Input:** Initial parameter $\{\theta^0, \Sigma_1^0, \dots, \Sigma_L^0\}$, training samples $\{x_i, y_i\}_{i=1}^n$, forecasting input x' , output model f , layer number L , repeating time in training m , repeating time in prediction k , epochs T , regularization $\{\alpha, \beta, \lambda\}$ and step size ξ .
- 2: **▷ Training process:**
for $t = 1, \dots, T$ **do**
 - 4: For each sample $\{x_i, y_i\}$, repeat Algorithm 2 m times, to receive its corresponding output $\{f_j^t(x_i)\}_{j=1}^m$.
Update parameter $\{\theta, \Sigma_1, \dots, \Sigma_L\}$ based on Eq (4) via gradient descent.
 - 6: **end for**
▷ Prediction process:
 - 8: Repeat Algorithm 2 k times with input x' , then receive its corresponding output $\{f_1(x'), \dots, f_k(x')\}$.
Calculate the final prediction based on:

$$f(x') = \frac{1}{k} \sum_{r=1}^k f_r(x').$$

10: **Output:** Final prediction $f(x')$.

Comparing with other methods. Injecting Gaussian noise is a widely used technique in data augmentation (Moreno-Barea et al., 2018) and adversarial training (Cohen et al., 2019) to improve generalization. In contrast, our proposed method goes beyond traditional augmentation in two key ways: (i) Rather than treating noise injection as a training-only strategy, we incorporate distributional projection directly into the model architecture, applying it consistently during both training and inference. This architectural integration enables theoretical guarantees for improved generalization. (ii) Instead of injecting noise at the input level only, we project the input into a learnable distribution at each layer, providing layerwise control over perturbations. This design offers better stability and helps prevent issues such as gradient explosion during training.

4.2 Theoretical Results

We provide some theoretical guarantees on why our approaches improve the Lipschitz and smoothness of the original model. Our analysis is focusing on one single neural network layer, which is denoted as $h(\cdot, \theta) : \mathbb{R}^p \rightarrow \mathbb{R}^q$:

$$h(\cdot, \theta) := [h_1(\cdot, \theta), \dots, h_q(\cdot, \theta)]^T.$$

Lipschitz norm. Johnson and Zhang (2023) proposed that there are two terms, i.e., “instability” term and “inconsistency” term, which are strongly predictive of the generalization gap (the difference between the performance on the training data and unseen data). Following Theorem 1 in Johnson and Zhang (2023), the “instability” term is related to the function Lipschitz. The following theorem demonstrates that our distributional input projection methods can reduce the function’s Lipschitz norm compared to the original model, contributing to improved generalization (The proof is in Appendix A.1).

Theorem 1. *Denote*

$$b_r := \max_x \|\nabla_x h_r(x, \theta)\|_2, \quad \mathcal{B}_r := \{x \in \mathbb{R}^p \mid \|\nabla_x h_r(x, \theta)\|_2 > c \cdot b_r\}, \quad r \in [q],$$

for some $0 < c < 1$. If the condition $\mu(\cup_r \mathcal{B}_r) < C$ holds, for any $\epsilon > 0$, there exists a distribution $\mathcal{P}$, such that

$$\max_{x \in \mathbb{R}^p} \|\mathbb{E}_{\eta \sim \mathcal{P}} \nabla_x h_r(x + \eta, \theta)\|_2 < (c + \epsilon) \cdot b_r, \quad \forall r \in [q].$$

Smoothness. It has been widely observed that lower function smoothness is often associated with better generalization performance in neural networks. Our proposed methods can also enforce the smoothness condition using a simple technique, as demonstrated in Theorem 2:

Theorem 2. *Denote*

$$s_r := \max_x \|\nabla_x^2 h_r(x, \theta)\|_2, \quad \mathcal{S}_r := \{x \in \mathbb{R}^p \mid \|\nabla_x^2 h_r(x, \theta)\|_2 > c \cdot s_r\}, \quad r \in [q],$$

for some $0 < c < 1$. If the condition $\mu(\cup_r \mathcal{S}_r) < C$ holds, for any $\epsilon > 0$, there exists a distribution $\mathcal{P}$, such that

$$\max_{x \in \mathbb{R}^p} \|\mathbb{E}_{\eta \sim \mathcal{P}} \nabla_x^2 h_r(x + \eta, \theta)\|_2 < (c + \epsilon) \cdot s_r, \quad \forall r \in [q].$$

The proof is in Appendix A.2. The results indicate that DIPNet can reduce the smoothness of models. In the sequel, we will show how this leads to better generalization performance.

Algorithm 2 DIPNet: Practical Implementation

- Input:** Model parameter $\{\theta, \Sigma_1, \dots, \Sigma_L\}$, input x , output model f , layer number L .
2: Initialize the input as v_0 .
for $l = 1, \dots, L$ **do**
4: Sample a particle u_l from $\mathcal{N}(v_{l-1}, \Sigma_l)$.
 Take u_l as the input of Layer- l .
6: Receive the output on Layer- l , and denote it as v_l .
end for
8: **Output:** v_L .
-

5 Experiment

We evaluate our proposed method against standard training and two common generalization methods—Sharpness-Aware Minimization (SAM) (Foret et al., 2020) and Randomized Smoothing (RS) (Cohen et al., 2019). Additional experiment explorations and ablation studies are in Appendix C.

5.1 Exploration on Vision Transformers under Adversarial Attacks

Compared to general non-perturbation methods, DIPNet achieves higher accuracy under adversarial attacks in training process, indicating their robustness. Here we consider two types of training-time attacks:

- **Randomized Gaussian noise:** we sample $\eta \sim \mathcal{N}(0, \sigma^2 I)$ for some $\sigma > 0$, and attack the input via $x \rightarrow x + \eta$ to simulate natural distribution shifts during training.
- **Fast Gradient Sign Method (FGSM) adversarial noise** (Goodfellow et al., 2014): we attack examples by $x \rightarrow x + \epsilon \cdot \text{sgn}(\nabla_x \mathcal{L}(\theta, x, y))$. This single-step attack efficiently approximates the worst-case attacks within an ℓ_∞ -ball of radius ϵ .

By injecting either type of attack into the training inputs, we can evaluate adversarial robustness by showing the accuracy on clean test data. The results show that DIPNet enhances stability under both worst-case attacks and random noise while preserving baseline performance. We first evaluate our method in Vision Transformers (ViTs) (Dosovitskiy et al., 2020), a popular architecture in computer vision known for its strong performance across classification, detection, and segmentation tasks.

Setup. We train three ViT variants—ViT-Tiny (5.5M parameters), ViT-Small (21.7M parameters), and ViT-Base (85.8M parameters)—on the CIFAR-100 dataset (Krizhevsky, 2009), an image classification dataset which contains 100 distinct classes of small-scale images. Each model is started from a checkpoint pretrained on ImageNet-21k and fine-tuned on ImageNet-1k, using a patch size of 16 and an input resolution of 224×224 . During training, we inject either additive Gaussian noise ($\sigma = 0.2$) or FGSM adversarial attacks ($\epsilon = 0.2$) into the training inputs. At evaluation, we report accuracy on clean test sets. Our implementation is adapted from the publicly available repository³, using default hyperparameters. Additional implementation details are provided in Appendix C.4.

Results. Table 1 shows the performance of various ViT models under different types of attack.

First, under the clean setting (no attack), our proposed method (DIPNet) consistently outperforms the Standard approach across all ViT variants (ViT-Tiny, ViT-Small, ViT-Base), demonstrating improved generalization capabilities. In contrast, RS, which introduces Gaussian perturbations at the input layer, significantly decreases clean accuracy, particularly evident in the ViT-Tiny model. See Appendix C.4 for training and validation curves.

Second, under Gaussian noise attacks, SAM and RS methods perform similarly to the Standard method, showing marginal improvements or even slight degradations. Conversely, our DIPNet method notably enhances robustness in the ViT-Tiny and ViT-Small and remains comparable to the baseline for the ViT-Base (further details on the ViT-Base Gaussian scenario are in Section 5.4).

Finally, under FGSM adversarial noise, our approach significantly outperforms both SAM and RS methods in the ViT-Tiny and ViT-Small. This highlights the effectiveness and superiority of DIPNet in mitigating adversarial attacks. However, on ViT-Base, RS achieves higher test accuracy. This likely occurs because RS’s randomized perturbation partially attenuates the FGSM attack—while DIPNet drives stronger adversarial fitting (higher train accuracy), RS can also weaken the attack and better prevent overfitting. See Appendix C.4 for training and validation curves.

Overall, our method consistently improves robustness across various attacks while preserving or slightly enhancing clean accuracy, indicating strong stability and generalization capabilities.

5.2 Exploration on LLM Reasoning

We extend our method to large-scale language models (LLMs), which often contain billions of parameters and now power a wide range of applications—from machine translation and summarization to

³<https://github.com/jeonsworld/ViT-pytorch>

Table 1: Test accuracy (%) on CIFAR-100 under different training-time attacks.

Attack type	Method	ViT-Tiny	ViT-Small	ViT-Base
None	Standard	84.71	89.59	92.46
	Sharpness-Aware Minimization (SAM)	84.46	89.72	92.68
	Randomized Smoothing (RS)	65.45	84.84	88.32
	DIPNet	85.36	90.06	92.84
Gaussian	Standard	46.31	75.65	69.13
	Sharpness-Aware Minimization (SAM)	46.04	74.22	69.28
	Randomized Smoothing (RS)	47.04	74.31	68.74
	DIPNet	51.62	78.21	69.31
FGSM	Standard	20.89	65.64	70.75
	Sharpness-Aware Minimization (SAM)	22.63	66.86	70.34
	Randomized Smoothing (RS)	23.22	65.55	77.30
	DIPNet	27.19	69.17	74.17

Table 2: Flexible-extract accuracy (%) on GSM8K.

Method	Qwen2.5-3B	Llama-3.2-3B	Gemma-3-4B-PT
Base Model	74.00	29.49	41.55
Standard Fine-tuning	74.45	32.68	42.00
Sharpness-Aware Minimization (SAM)	74.22	31.84	39.65
Randomized Smoothing (RS)	75.59	33.06	41.24
DIPNet	75.66	35.56	42.08

dialogue systems and deep thinking. In particular, mathematical reasoning tasks like GSM8K (Cobbe et al., 2021) have recently attracted significant attention for assessing an LLM’s ability to perform precise arithmetic and logical reasoning. We hypothesize that our method will help LLMs form smoother latent representations, improving both generalization to new problems and robustness against prompt variations.

Setup. We use the official OpenAI GSM8K dataset, which consists of 1,319 grade-school math word problems requiring several reasoning steps and an exact numerical answer, replacing the original “#### <answer>” marker with “The answer is <answer>” and removing all calculator annotations (e.g., “<<200×5=1000>>”) to match the LM Evaluation Harness (Gao et al., 2024)’s answer-extraction format. To test our method across diverse architectures, we select three popular open-source models: Qwen2.5-3B (Yang et al., 2024), Llama-3.2-3B (Dubey et al., 2024), and Gemma-3-4B-PT (Farabet and Warkentin, 2025). Due to computational constraints, we fine-tune each model with LoRA for a single epoch—using a learning rate of 10^{-4} for Qwen2.5-3B and 2×10^{-4} for the other two—then evaluate with the LM Evaluation Harness on the GSM8K CoT task. Additional implementation details are provided in Appendix C.5.

Results. Table 2 reports flexible-extract accuracy (any correct numeric output). Our method consistently outperforms other baselines on all three models, yielding nontrivial gains. Results for strict-match accuracy are given in Appendix C.5, where we observe even larger improvements in Qwen2.5-3B and Llama-3.2-3B. These findings suggest that our method not only improve numeric correctness but also enhance the model’s ability to generalize to answer formats.

5.3 Out-Of-Distribution (OOD) Evaluation

We also explore the generalization performance on OOD tasks. We conduct experiments on tabular datasets using different neural network architectures, further validating the advantages of our proposed methods in both In-Distribution (ID) and Out-Of-Distribution (OOD) generalization performance.

Setup. To assess robustness under distribution shift, we use Shifts Weather Prediction dataset (Malinin et al., 2021), which provides a scalar regression task to forecast the air temperature. We train three MLP architectures—2 layers with 100 neurons per layer (2+100), 4 layers with 100 neurons per




Figure 2: ID and OOD MSE on Weather Dataset.

Notes. Here we train on the ID dataset, validate on both ID and OOD datasets, and test the model performance on both ID and OOD datasets (with the test OOD environment differing from the validation OOD environment).

layer (4+100), and 4 layers with 400 neurons per layer (4+400)—using standard stochastic gradient descent on the MSE loss. We compare against the non-perturbed baseline (Standard). To study performance under different data regimes, we vary the fraction of available training data (100%, 66%, and 33%) and evaluate all three architectures under each split.

Results. Figure 2 presents both ID and OOD test MSEs. The results shows that comparing with standard non-perturbation method, DIPNet consistently improve ID performance, and enhance OOD generalization significantly. More detailed results are shown in Appendix C.2.

5.4 Noise-Level Ablation: Gaussian Corruption Effects Across Model Sizes

To quantify the impact of our DIPNet module under varying gaussian noise levels, we evaluate both ViT-Base and ViT-Tiny corrupted by additive Gaussian noise with standard deviation $\sigma \in \{0.2, 0.5, 1.0\}$. Table 3 reports test accuracy for the standard (baseline) training and for our method.

For larger model, DIPNet yields consistent gains over the baseline at all noise levels. The gap widens as noise increases, demonstrating that DIPNet effectively enhances robustness to heavy corruption. For smaller model, while DIPNet brings large relative gains at low to moderate noise, the absolute accuracy remains low. This indicates that ViT-Tiny’s limited capacity leads to underfitting when faced with strong noise, and suggests that further architectural capacity or noise-specific regularization may be required. These ablation results confirm that DIPNet consistently improves noise robustness, but also highlight the interaction between module efficacy and backbone capacity under severe corruption.

Table 3: Evaluation under Gaussian Noise Levels ($\sigma \in \{0.2, 0.5, 1.0\}$).

Model	Method	0.2	0.5	1.0
ViT-Base	Standard	69.13	54.81	35.02
	DIPNet	69.31	54.95	36.08
ViT-Tiny	Standard	46.31	26.99	15.93
	DIPNet	51.62	32.75	17.63

6 Conclusion

In this work, we proposed Distributional Input Projection Networks (DIPNet), a novel architectural framework designed to enhance generalization by projecting inputs into learnable distributions at each layer. This approach implicitly enforces smoothness and reduces the Lipschitz constant of the network, supported by both theoretical analysis and extensive empirical validation. Across a wide range of models and tasks, DIPNet consistently improves test performance in standard, adversarial, and out-of-distribution settings. Our results demonstrate that DIPNet offers a broadly applicable and effective strategy for improving generalization in overparameterized deep networks. We consider extending this framework to reasoning tasks and reinforcement learning as a promising direction for future work.

Acknowledgements

We would like to thank Rui Pan for his helpful comments and suggestions.

References

- Abbas, M., Xiao, Q., Chen, L., Chen, P.-Y., and Chen, T. (2022). Sharp-maml: Sharpness-aware model-agnostic meta learning. In *International conference on machine learning*, pages 10–32. PMLR.
- Bahri, D., Mobahi, H., and Tay, Y. (2021). Sharpness-aware minimization improves language model generalization. *arXiv preprint arXiv:2110.08529*.
- Bousquet, O. and Elisseeff, A. (2002). Stability and generalization. *The Journal of Machine Learning Research*, 2:499–526.
- Chen, X., Hsieh, C.-J., and Gong, B. (2021). When vision transformers outperform resnets without pre-training or strong data augmentations. *arXiv preprint arXiv:2106.01548*.
- Cobbe, K., Kosaraju, V., Bavarian, M., Chen, M., Jun, H., Kaiser, L., Plappert, M., Tworek, J., Hilton, J., Nakano, R., et al. (2021). Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*.
- Cohen, J., Rosenfeld, E., and Kolter, Z. (2019). Certified adversarial robustness via randomized smoothing. In *international conference on machine learning*, pages 1310–1320. PMLR.
- Dhariwal, P. and Nichol, A. (2021). Diffusion models beat gans on image synthesis. *Advances in neural information processing systems*, 34:8780–8794.
- Dobriban, E., Hassani, H., Hong, D., and Robey, A. (2023). Provable tradeoffs in adversarially robust classification. *IEEE Transactions on Information Theory*.
- Donhauser, K., Tifrea, A., Aerni, M., Heckel, R., and Yang, F. (2021). Interpolation can hurt robust generalization even when there is no noise. *Advances in Neural Information Processing Systems*, 34:23465–23477.
- Dosovitskiy, A., Beyer, L., Kolesnikov, A., Weissenborn, D., Zhai, X., Unterthiner, T., Dehghani, M., Minderer, M., Heigold, G., Gelly, S., Uszkoreit, J., and Houlsby, N. (2020). An image is worth 16x16 words: Transformers for image recognition at scale. *arXiv preprint arXiv:2010.11929*.
- Dubey, A., Jauhri, A., Pandey, A., Kadian, A., Al-Dahle, A., Letman, A., Mathur, A., Schelten, A., Yang, A., Fan, A., et al. (2024). The llama 3 herd of models. *arXiv preprint arXiv:2407.21783*.
- Fanaee-T, H. and Gama, J. (2014). Event labeling combining ensemble detectors and background knowledge. *Progress in Artificial Intelligence*, 2:113–127.
- Farabet, C. and Warkentin, T. (2025). Introducing gemma 3: The most capable model you can run on a single gpu or tpu. <https://blog.google/technology/developers/gemma-3/>.
- Foret, P., Kleiner, A., Mobahi, H., and Neyshabur, B. (2020). Sharpness-aware minimization for efficiently improving generalization. *arXiv preprint arXiv:2010.01412*.
- Gao, L., Tow, J., Abbasi, B., Biderman, S., Black, S., DiPofi, A., Foster, C., Golding, L., Hsu, J., Le Noac’h, A., Li, H., McDonell, K., Muennighoff, N., Ociepa, C., Phang, J., Reynolds, L., Schoelkopf, H., Skowron, A., Sutawika, L., Tang, E., Thite, A., Wang, B., Wang, K., and Zou, A. (2024). The language model evaluation harness.
- Goodfellow, I. J., Shlens, J., and Szegedy, C. (2014). Explaining and harnessing adversarial examples. *arXiv preprint arXiv:1412.6572*.
- Gowal, S., Qin, C., Uesato, J., Mann, T., and Kohli, P. (2020). Uncovering the limits of adversarial training against norm-bounded adversarial examples. *arXiv preprint arXiv:2010.03593*.

- Hao, Y. and Zhang, T. (2024). The surprising harmfulness of benign overfitting for adversarial robustness. *arXiv preprint arXiv:2401.12236*.
- Ho, J., Jain, A., and Abbeel, P. (2020). Denoising diffusion probabilistic models. *Advances in neural information processing systems*, 33:6840–6851.
- Hochreiter, S. and Schmidhuber, J. (1997). Flat minima. *Neural computation*, 9(1):1–42.
- Izmailov, P., Podoprikin, D., Garipov, T., Vetrov, D., and Wilson, A. G. (2018). Averaging weights leads to wider optima and better generalization. *arXiv preprint arXiv:1803.05407*.
- Javanmard, A., Soltanolkotabi, M., and Hassani, H. (2020). Precise tradeoffs in adversarial training for linear regression. In *Conference on Learning Theory*, pages 2034–2078. PMLR.
- Jiang, Y., Nagarajan, V., Baek, C., and Kolter, J. Z. (2021). Assessing generalization of sgd via disagreement. *arXiv preprint arXiv:2106.13799*.
- Jiang, Y., Neyshabur, B., Mobahi, H., Krishnan, D., and Bengio, S. (2019). Fantastic generalization measures and where to find them. *arXiv preprint arXiv:1912.02178*.
- Johnson, R. and Zhang, T. (2023). Inconsistency, instability, and generalization gap of deep neural network training. *arXiv preprint arXiv:2306.00169*.
- Keskar, N. S., Mudigere, D., Nocedal, J., Smelyanskiy, M., and Tang, P. T. P. (2016). On large-batch training for deep learning: Generalization gap and sharp minima. *arXiv preprint arXiv:1609.04836*.
- Kirsch, A. and Gal, Y. (2022). A note on "assessing generalization of sgd via disagreement". *arXiv preprint arXiv:2202.01851*.
- Krizhevsky, A. (2009). Learning multiple layers of features from tiny images. Technical report, University of Toronto. Technical Report.
- Lecuyer, M., Atlidakis, V., Geambasu, R., Hsu, D., and Jana, S. (2019). Certified robustness to adversarial examples with differential privacy. In *2019 IEEE symposium on security and privacy (SP)*, pages 656–672. IEEE.
- Madry, A., Makelov, A., Schmidt, L., Tsipras, D., and Vladu, A. (2017). Towards deep learning models resistant to adversarial attacks. *arXiv preprint arXiv:1706.06083*.
- Malinin, A., Athanasopoulos, A., Barakovic, M., Cuadra, M. B., Gales, M. J., Granziera, C., Graziani, M., Kartashev, N., Kyriakopoulos, K., Lu, P.-J., et al. (2022). Shifts 2.0: Extending the dataset of real distributional shifts. *arXiv preprint arXiv:2206.15407*.
- Malinin, A., Band, N., Chesnokov, G., Gal, Y., Gales, M. J., Noskov, A., Ploskonosov, A., Prokhorenkova, L., Provilkov, I., Raina, V., et al. (2021). Shifts: A dataset of real distributional shift across multiple large-scale tasks. *arXiv preprint arXiv:2107.07455*.
- Moreno-Barea, F. J., Strazzera, F., Jerez, J. M., Urda, D., and Franco, L. (2018). Forward noise adjustment scheme for data augmentation. In *2018 IEEE symposium series on computational intelligence (SSCI)*, pages 728–734. IEEE.
- Nakkiran, P. and Bansal, Y. (2020). Distributional generalization: A new kind of generalization. *arXiv preprint arXiv:2009.08092*.
- Qin, Y., Song, D., Chen, H., Cheng, W., Jiang, G., and Cottrell, G. (2017). A dual-stage attention-based recurrent neural network for time series prediction. *arXiv preprint arXiv:1704.02971*.
- Rombach, R., Blattmann, A., Lorenz, D., Esser, P., and Ommer, B. (2022). High-resolution image synthesis with latent diffusion models. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 10684–10695.
- Saharia, C., Chan, W., Chang, H., Lee, C., Ho, J., Salimans, T., Fleet, D., and Norouzi, M. (2022). Palette: Image-to-image diffusion models. In *ACM SIGGRAPH 2022 conference proceedings*, pages 1–10.

- Salman, H., Li, J., Razenshteyn, I., Zhang, P., Zhang, H., Bubeck, S., and Yang, G. (2019). Provably robust deep learning via adversarially trained smoothed classifiers. *Advances in neural information processing systems*, 32.
- Shalev-Shwartz, S., Shamir, O., Srebro, N., and Sridharan, K. (2010). Learnability, stability and uniform convergence. *The Journal of Machine Learning Research*, 11:2635–2670.
- Shen, X. and Meinshausen, N. (2023). Engression: Extrapolation for nonlinear regression? *arXiv preprint arXiv:2307.00835*.
- Song, J., Meng, C., and Ermon, S. (2020). Denoising diffusion implicit models. *arXiv preprint arXiv:2010.02502*.
- Tsipras, D., Santurkar, S., Engstrom, L., Turner, A., and Madry, A. (2018). Robustness may be at odds with accuracy. *arXiv preprint arXiv:1805.12152*.
- Vito, S. (2016). Air Quality. UCI Machine Learning Repository. DOI: <https://doi.org/10.24432/C59K5F>.
- Wang, Y., Ma, X., Bailey, J., Yi, J., Zhou, B., and Gu, Q. (2021). On the convergence and robustness of adversarial training. *arXiv preprint arXiv:2112.08304*.
- Yang, A., Yang, B., Zhang, B., Hui, B., Zheng, B., Yu, B., Li, C., Liu, D., Huang, F., Wei, H., et al. (2024). Qwen2. 5 technical report. *arXiv preprint arXiv:2412.15115*.
- Yang, G., Duan, T., Hu, J. E., Salman, H., Razenshteyn, I., and Li, J. (2020). Randomized smoothing of all shapes and sizes. In *International Conference on Machine Learning*, pages 10693–10705. PMLR.
- Yue, Y., Jiang, J., Ye, Z., Gao, N., Liu, Y., and Zhang, K. (2023). Sharpness-aware minimization revisited: Weighted sharpness as a regularization term. In *Proceedings of the 29th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, pages 3185–3194.
- Zhang, H., Yu, Y., Jiao, J., Xing, E., El Ghaoui, L., and Jordan, M. (2019). Theoretically principled trade-off between robustness and accuracy. In *International conference on machine learning*, pages 7472–7482. PMLR.
- Zhou, H., Zhang, S., Peng, J., Zhang, S., Li, J., Xiong, H., and Zhang, W. (2021). Informer: Beyond efficient transformer for long sequence time-series forecasting. In *The Thirty-Fifth AAAI Conference on Artificial Intelligence, AAAI 2021, Virtual Conference*, volume 35, pages 11106–11115. AAAI Press.
- Zou, D., Frei, S., and Gu, Q. (2021). Provable robustness of adversarial training for learning halfspaces with noise. In *International Conference on Machine Learning*, pages 13002–13011. PMLR.

A Theoretical Analysis

A.1 Proof of Theorem 1

We can consider $\mathcal{P}$ as a uniform distribution on $\mathbb{R}^p$, which is supported on a set with measurement ζC , then for any $r \in [q]$, we can obtain that

$$\begin{aligned} \|\mathbb{E}_\eta \nabla_x h_r(x + \eta, \theta)\|_2^2 &\leq \mathbb{E}_\eta \|\nabla_x h_r(x + \eta, \theta)\|_2^2 \\ &= \mathbb{E}_\eta [\|\nabla_x h_r(x + \eta, \theta)\|_2^2 | 1(x + \eta \in \cup_r \mathcal{B}_r)] \mathbb{P}(\cup_r \mathcal{B}_r) \\ &\quad + \mathbb{E}_\eta [\|\nabla_x h_r(x + \eta, \theta)\|_2^2 | 1(x + \eta \notin \cup_r \mathcal{B}_r)] \mathbb{P}((\cup_r \mathcal{B}_r)^c) \\ &\leq \frac{1}{\zeta C} (C \cdot b_r + \zeta C \cdot cb_r) = cb_r + \frac{1}{\zeta} b_r. \end{aligned}$$

Choosing $\zeta > \epsilon^{-1}$, we can finish the proof.

A.2 Proof of Theorem 2

The proof is similar to the proof for Theorem 1. We can consider $\mathcal{P}$ as a uniform distribution on $\mathbb{R}^p$, which is supported on a set with measurement ζC , then for any $r \in [q]$, we can obtain that

$$\begin{aligned} \|\mathbb{E}_\eta \nabla_x^2 h_r(x + \eta, \theta)\|_2^2 &\leq \mathbb{E}_\eta \|\nabla_x^2 h_r(x + \eta, \theta)\|_2^2 \\ &= \mathbb{E}_\eta [\|\nabla_x^2 h_r(x + \eta, \theta)\|_2^2 | 1(x + \eta \in \cup_r \mathcal{S}_r)] \mathbb{P}(\cup_r \mathcal{S}_r) \\ &\quad + \mathbb{E}_\eta [\|\nabla_x^2 h_r(x + \eta, \theta)\|_2^2 | 1(x + \eta \notin \cup_r \mathcal{S}_r)] \mathbb{P}((\cup_r \mathcal{S}_r)^c) \\ &\leq \frac{1}{\zeta C} (C \cdot s_r + \zeta C \cdot cs_r) = cs_r + \frac{1}{\zeta} s_r. \end{aligned}$$

Choosing $\zeta > \epsilon^{-1}$, we can finish the proof.

B Dataset Description

B.1 Tabular Dataset for MLP

- **Air Quality.** (Vito, 2016) The Air Quality dataset comprises 15 features derived from hourly averaged responses of an array of 5 metal oxide chemical sensors embedded in an Air Quality Chemical Multisensor Device in an Italian city, collected from March 2004 to February 2005. In this study, we use 9 of these features, with “PT08.S3(NOx)” as the response variable and the remaining 8 features as input variables.
- **ETD.** (Zhou et al., 2021) The ETD dataset is related to the electronic data distribution hourly data recordings. It contains 8 features, including the date of the point, 6 different types of external power load features and the predictive value “oil temperature”.
- **Sharing Bike.** (Fanaee-T and Gama, 2014) This dataset aims to understand the factors influencing the demand for shared bikes in the American market, with hourly data recorded from January 2011 to December 2012. For this analysis, we consider 4 features, i.e., weather, temperature, humidity, and feeling temperature, to predict the total count of rental bikes.
- **NASDAQ100.** (Qin et al., 2017) It contains stock prices of 81 corporations under NASDAQ 100 and the index value of NASDAQ 100. The frequency of the data collection is one-minute. Here we consider the stock prices of different corporations as input variables, and the index value of NASDAQ 100 as predictive variable.

Table 4: Overview of tabular dataset.

Dataset	Air Quality	ETD	Sharing Bike	NASDAQ100
input task	vector	vector	vector	vector
# samples	8991	17420	17379	40560


Table 5: Overview of OOD tabular dataset.

Dataset	$VPower_s$	$VPower_r$	Weather
input task	vector	vector	vector
# samples	546543	554642	436733

B.2 OOD Dataset

To assess robustness under distribution shift, we use the Shifts vessel power estimation dataset (Malinin et al., 2022) and the Shifts Weather Prediction dataset (Malinin et al., 2021) in OOD situations (see details in Table 5 and Figure 3).

- **Vessel Power Estimation**⁴ The Shifts Vessel Power Prediction dataset contains 11 different features and one scalar target (energy utilization of cargo ships). And it contains two types of data: one is synthetic dataset ($VPower_s$), the other is real dataset ($VPower_r$).
- **Weather Prediction**⁵ The Shifts Weather Prediction dataset (Malinin et al., 2021) provides a scalar regression task to forecast the air temperature. The dataset contains 127 features and a single regression target, spanning an entire year, i.e., from September 1st, 2018, to September 1st, 2019.



Notes. The figure on the left side is about the Shifts vessel power estimation dataset, and the figure on the right is about the Shifts Weather Prediction dataset. “train”, “dev” and “eval” refer to training data, validation data and test data respectively. As public access to canonical “test” of vessel power dataset is restricted, here we valid on “dev-in” and use “dev-out” to test the model performance.

Figure 3: OOD dataset descriptions

C Supplement Experiments

C.1 Explorations on MLP

Here we compare the generalization performance of our method and a standard non-perturbation output model by evaluating them on various datasets and different neural network architectures.

To measure the generalization performance of different methods, we take explorations on four tabular datasets and detailed description could be found in Appendix B.

Specifically, we test different settings, including various proportions of training and test data, the number of hidden layers, and the number of neurons in each layer. Using gradient descent on the ℓ_2 loss, we summarize the results in Table 6, indicating that both the standard errors and the adversarial errors on DIPNet are consistently smaller than those for standard networks. This further verifies the benefits of our proposed method.

C.2 Details for Out-Of-Distribution (OOD) Evaluation

Here we explore the generalization performance on Out-Of-Distribution (OOD) tasks. We conduct experiments on three tabular datasets using different neural network architectures, further validating the advantages of our proposed methods in OOD generalization performance.

Setup. To assess robustness under distribution shift, we use the Shifts vessel power estimation dataset (Malinin et al., 2022) and the Shifts Weather Prediction dataset (Malinin et al., 2021) in OOD situations. For each dataset, we train three MLP architectures—2 layers with 100 neurons per layer (2+100), 4 layers with 100 neurons per layer (4+100), and 4 layers with 400 neurons per

⁴Notice that this is a developing dataset and it only contains the training data and valid data recently. So here we use the ID valid dataset to take validation, and the OOD valid dataset is used for testing.

⁵More details can be found in <https://github.com/Shifts-Project/shifts>.

Table 6: Standard and Adversarial MSE on MLP.

Dataset	Test proposition	Method	St(2+100)	Adv(2+100)	St(4+100)	Adv(4+100)	St(4+400)	Adv(4+400)
Air Quality	0.3	Standard	0.0858	24.0112	0.0677	37.1026	0.0628	20.5767
		DIPNet	0.0822	11.14	0.0576	20.9601	0.0526	20.5601
	0.5	Standard	0.1033	13.8045	0.0873	49.1966	0.0823	17.7776
		DIPNet	0.0964	9.6509	0.0794	5.8675	0.0748	12.9039
	0.7	Standard	0.0984	16.6161	0.0938	24.4145	0.0910	14.5242
		DIPNet	0.0901	4.8616	0.0844	6.6429	0.0853	10.1404
ETD	0.3	Standard	0.1845	4.7174	0.1693	5.1239	0.1635	4.5290
		DIPNet	0.1787	3.6371	0.1604	3.7268	0.1578	2.9117
	0.5	Standard	0.1905	6.1011	0.1724	4.5385	0.1668	4.5335
		DIPNet	0.1831	3.5267	0.1648	3.3667	0.1632	3.1108
	0.7	Standard	0.2023	5.0696	0.1913	6.2199	0.1867	5.1136
		DIPNet	0.1947	4.3076	0.1776	2.9721	0.1757	2.9267
Sharing Bike	0.3	Standard	0.7221	2.5177	0.7208	2.0264	0.7216	2.1557
		DIPNet	0.7188	2.0034	0.7186	1.8647	0.7182	1.8732
	0.5	Standard	0.7108	2.3387	0.7107	2.3030	0.7118	1.8314
		DIPNet	0.7084	1.7799	0.7082	1.8780	0.7089	1.7272
	0.7	Standard	0.7075	1.7905	0.7057	1.8397	0.7068	1.7931
		DIPNet	0.7050	1.7353	0.7043	1.7941	0.7056	1.7785
NASDAQ100	0.3	Standard	3.576e-5	0.6066	3.538e-5	0.5447	2.779e-5	0.5303
		DIPNet	2.757e-5	0.5824	2.511e-5	0.5628	2.158e-5	0.5195
	0.5	Standard	3.691e-5	0.5901	3.671e-5	0.5364	2.916e-5	0.5389
		DIPNet	2.890e-5	0.5866	2.642e-5	0.5805	2.266e-5	0.5244
	0.7	Standard	6.146e-5	0.5813	8.085e-5	0.5506	3.612e-5	0.5455
		DIPNet	3.122e-5	0.5812	2.886e-5	0.5428	2.532e-5	0.5336

Notes. “St”, “Adv” refer to standard loss and adversarial loss respectively. “2+100”, “4+100”, “4+400” refer to the network architecture (e.g. “2+100” means the two layer network which each layer contains 100 neurons).

layer (4+400)—using standard stochastic gradient descent on the MSE loss. We compare against the non-perturbed baseline (Standard). We vary the fraction of training data (100%, 66%, 33%) and evaluate the same three MLP architectures under each split.

Results. Tables 7 and 8 present OOD test MSEs. The OOD test error results show that DIPNet consistently achieve better OOD generalization performance compared to the non-perturbation method, indicating that input perturbations enhance OOD generalization.

Table 7: OOD MSE on Vessel Power dataset.

% Train		100%			66%			33%		
Dataset	Method	2+100	4+100	4+400	2+100	4+100	4+400	2+100	4+100	4+400
V_{power_s}	Standard	0.0212	0.0231	0.0288	0.0401	0.0533	0.0601	0.0539	0.0498	0.0420
	DIPNet	0.0195	0.0206	0.0212	0.0232	0.0305	0.0298	0.0298	0.0318	0.0324
V_{power_r}	Standard	0.0400	0.0560	0.0895	0.0439	0.0401	0.0730	0.0944	0.0947	0.0859
	DIPNet	0.0363	0.0439	0.0555	0.0345	0.0383	0.0412	0.0581	0.0534	0.0593

Notes. “2+100”, “4+100”, “4+400” refer to the network architecture (e.g. “2+100” means the two layer network which each layer contains 100 neurons).

Table 8: ID and OOD MSE on Weather dataset.

% Train	Method	ID(2+100)	ID(4+100)	ID(4+400)	OOD(2+100)	OOD(4+100)	OOD(4+400)
100 %	Standard	0.0360	0.0330	0.0291	0.0640	0.0647	0.0576
	DIPNet	0.0344	0.0322	0.0280	0.0562	0.0528	0.0547
66 %	Standard	0.0364	0.0329	0.0315	0.0898	0.0723	0.0822
	DIPNet	0.0354	0.0324	0.0295	0.0558	0.0552	0.0569
33 %	Standard	0.0364	0.0341	0.0325	0.0741	0.0599	0.0662
	DIPNet	0.0361	0.0337	0.0320	0.0519	0.0549	0.0627

Notes. Here we train on the ID dataset, validate on both ID and OOD datasets, and test the model performance on both ID and OOD datasets (with the test OOD environment differing from the validation OOD environment). “2+100”, “4+100”, “4+400” refer to the network architecture (e.g. “2+100” means the two layer network which each layer contains 100 neurons). “ID” means test data on ID environment, and “OOD” means test data on OOD environment.

C.3 Explorations on ResNet

Furthermore, we extend our methods to ResNet architecture. To examine its performance compared to the standard non-perturbation ResNet, we conducted experiments on CIFAR10 and CIFAR100 using different ResNet structures. We then compared the classification accuracy (in percentage) on the test data for each method. The results, summarized in Table 9, demonstrate the benefits of DIPNet in improving generalization performance. These findings highlighting the extended potential of our proposed methods.

Table 9: Result classification accuracy (%) on ResNet.

Method	res18-cifar10	res34-cifar10	res50-cifar10	wideres-cifar10	wideres-cifar100
Standard	90.96	93.52	93.37	96.14	80.77
DIPNet	91.72	93.56	93.89	96.53	81.73

C.4 Details for ViT Experiments

All Vision Transformer experiments are conducted on a single NVIDIA A100-80G GPU. Checkpoints are loaded via the `timm` library. We enable Apex O2 mixed-precision training (FP16) with all other hyperparameters at their defaults.

For all baselines and our method, we only sample once at evaluation.

For the two baselines—SAM and RS—we first perform a hyperparameter sweep on the ViT-Tiny model, then scale the optimal settings to larger backbones. Specifically, SAM’s perturbation radius ρ is searched over $\{0.01, 0.03, 0.05\}$, yielding $\rho = 0.05$; RS’s noise level σ is searched over $\{0.001, 0.005, 0.01, 0.05, 0.1\}$, yielding $\sigma = 0.01$.

Representative training and validation curves can be found in Figure 4, 5, 6 and 7.

C.5 Details for LLM Experiments

All language-model experiments are conducted on a single NVIDIA RTX 4090 GPU. We apply LoRA with rank 8, alpha 16, dropout 0.1, targeting the modules ["q_proj", "v_proj"].

Training arguments include a batch size of 64, 1 epoch, weight decay of 0.01, warmup ratio of 0.03, and a fixed seed of 42. We search peak learning rates in $\{7 \times 10^{-5}, 1 \times 10^{-4}, 2 \times 10^{-4}, 3 \times 10^{-4}, 5 \times 10^{-4}\}$, which yields 1×10^{-4} for Qwen2.5-1.5B and 2×10^{-4} for both Llama-3.2-3B and Gemma-3-4B-PT.

Evaluation is performed on the GSM8K CoT task with 8 shots, a batch size of 52, and seed 42.

For the two baselines—SAM and RS—we conduct hyperparameter sweeps and report the best-performing accuracy. In this case, SAM’s ρ is searched over $\{0.01, 0.02, 0.05\}$ and RS’s σ over $\{0.01, 0.02, 0.05\}$.

Table 10: Strict-match accuracy (%) on GSM8K.

Method	Qwen2.5-3B	Llama-3.2-3B	Gemma-3-4B-PT
Base Model	58.83	27.45	39.95
Standard Fine-tuning	68.61	31.92	41.55
Sharpness-Aware Minimization (SAM)	69.37	30.86	39.27
Randomized Smoothing (RS)	69.45	32.22	40.79
DIPNet	71.57	35.41	41.55

C.6 Additional Ablation Study

C.6.1 Ablation Study on Layerwise Distributional Input Projection

To assess how the depth of distributional input projections affects robustness, we compare four configurations of Layerwise DIPNet: taking the distributional input projection only at the first layer

Table 11: **Ablation studies on ViT under Gaussian attack.** This table reports test accuracy (%) on the CIFAR-100 dataset. “Standard” is the baseline ViT without any distributional input projection; “Full-Layer” (in a) and “Learnable” (in b) are the same learnable full distributional projection-at-every-layer setting.

(a) Depth Ablation (Layerwise DIPNet).

Method	ViT-Tiny	ViT-Small	ViT-Base
Standard	46.31	75.65	69.13
1-Layer	50.39	75.55	68.81
2-Layer	51.90	75.99	71.78
Full-Layer	51.62	78.21	69.31

(b) Perturbation Coefficient Ablation (Learnable vs. Fixed DIPNet).

Method	ViT-Tiny	ViT-Small	ViT-Base
Standard	46.31	75.65	69.13
Fixed-0.5	51.91	77.02	69.95
Fixed-1.0	48.94	75.97	67.23
Fixed-1.2	47.44	75.24	65.93
Fixed-Learned	45.24	72.29	64.82
Learnable	51.62	78.21	69.31

(“1-Layer”), at the first two layers (“2-Layer”), at every layer (“Full-Layer”), and the original model without any distributional projection (“Standard”).


As Table 11a shows, taking the distributional input projection at only a single layer (“1-Layer”) yields little to no generalization benefit—and can even lead to a slight drop in accuracy. In contrast, the “2-Layer” configuration delivers a clear generalization boost, outperforming Standard across all ViT scales, and on both ViT-Tiny and ViT-Base it even edges out the “Full-Layer” configuration—most notably pushing ViT-Base to 71.78 %. Overall, taking the distributional input projection at multiple depths delivers substantial gains, although the optimal number of layers can vary by model scale.

C.6.2 Ablation Study on Learnable Distributional Input Projection


We then explore different strategies for the learnable parameter `coef`—which controls the magnitude of the distributional projection added at each layer—in DIPNet: fixing `coef` to a constant (0.5, 1.0 or 1.2) or the optimal value obtained from a learnable run (“Fixed-Learned”), or treating `coef` as a trainable parameter initialized randomly (“Learnable”). In the “Fixed-Learned” setting, we initialize `coef` to the value obtained from “Learnable” setting (1.4121 for ViT-Tiny, -1.4141 for ViT-Small, and 1.4141 for ViT-Base). As Table 11b shows, fixing `coef` to a small constant produces an accuracy comparable to that of the “Learnable” setting, even though the “Learnable” setting itself starts from a random small value and grows to a larger optimum. However, choosing a larger fixed `coef` leads to a clear drop in performance, even when that constant remains below the final magnitude learned by the “Learnable” setting. This contrast highlights the advantage of a learnable `coef`, which can flexibly adjust perturbation strength to the ideal level for all input layers.

D Impact statement


This paper contributes foundational research in the areas of model generalization within the machine learning community. Our primary goal is to advance the practical methodologies to improve model generalization. Given the scope of this research, we do not anticipate immediate ethical concerns or direct societal consequences. Therefore, we believe there are no specific ethical considerations or immediate societal impacts to be emphasized in the context of this work.




(a) Training Loss



(b) Validation Loss




(c) Training Accuracy




(d) Validation Accuracy


Figure 4: Training and validation curves on ViT-Tiny without any attacks.




(a) Training Loss



(b) Validation Loss




(c) Training Accuracy




(d) Validation Accuracy


Figure 5: Training and validation curves on ViT-Base under FGSM adversarial attacks.




(a) Training Loss



(b) Validation Loss



(c) Training Accuracy




(d) Validation Accuracy


Figure 6: Training and validation curves on ViT-Tiny under Gaussian noise.




(a) Training Loss



(b) Validation Loss



(c) Training Accuracy



(d) Validation Accuracy

Figure 7: Training and validation curves on ViT-Tiny under FGSM adversarial attacks.